



OpenClassRooms

Data Scientist

P8 Deploy a model in the cloud



Fruits!

Pictures used for educational purpose only

Sommaire

I. Introduction

II. Big Data environment

III. Image processing with Spark


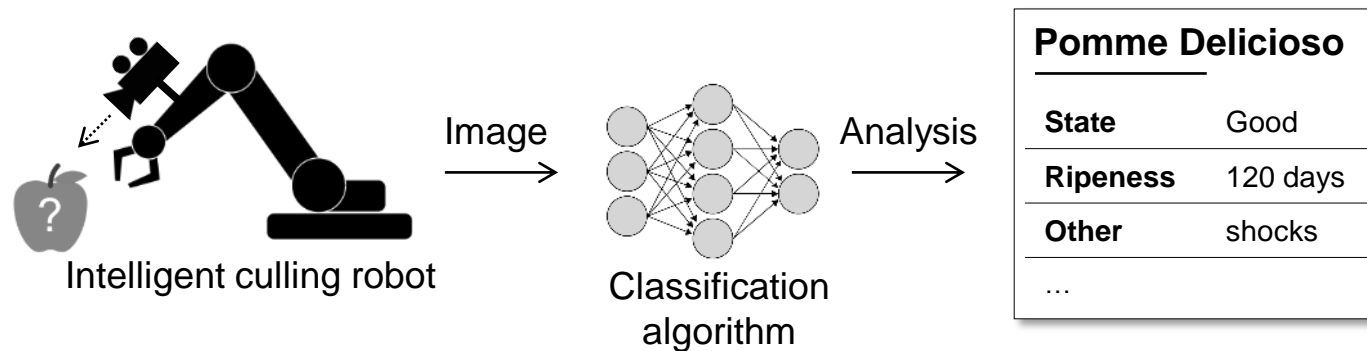
IV. Conclusion

I. Introduction

1) The project


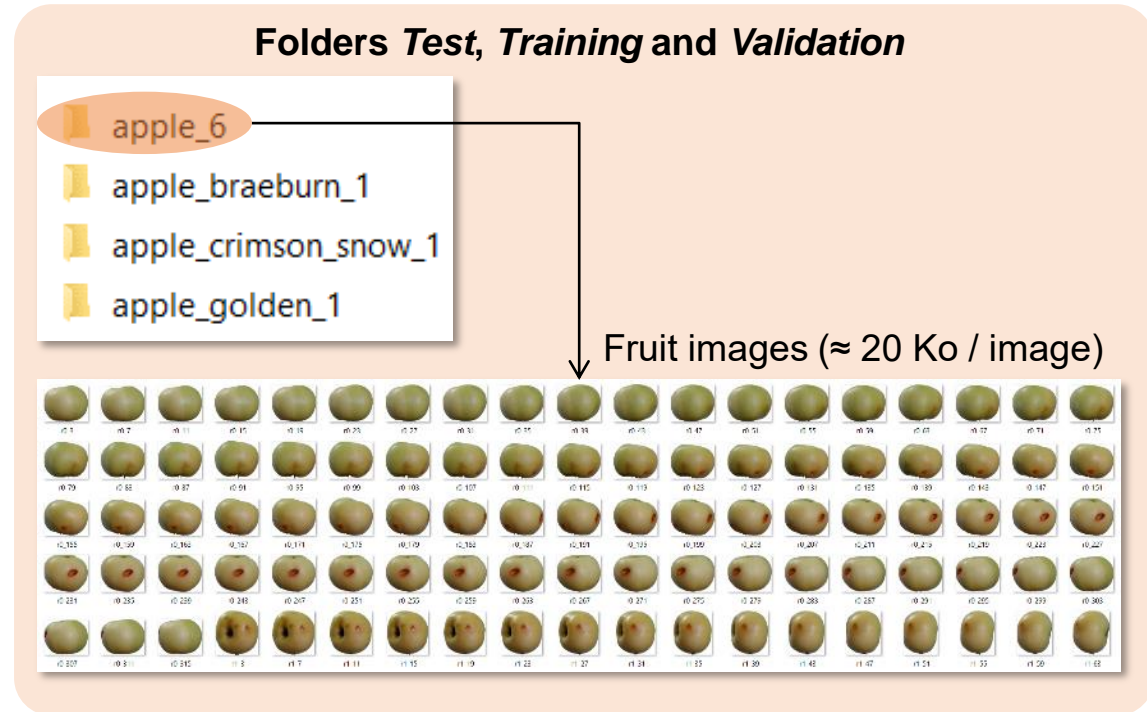
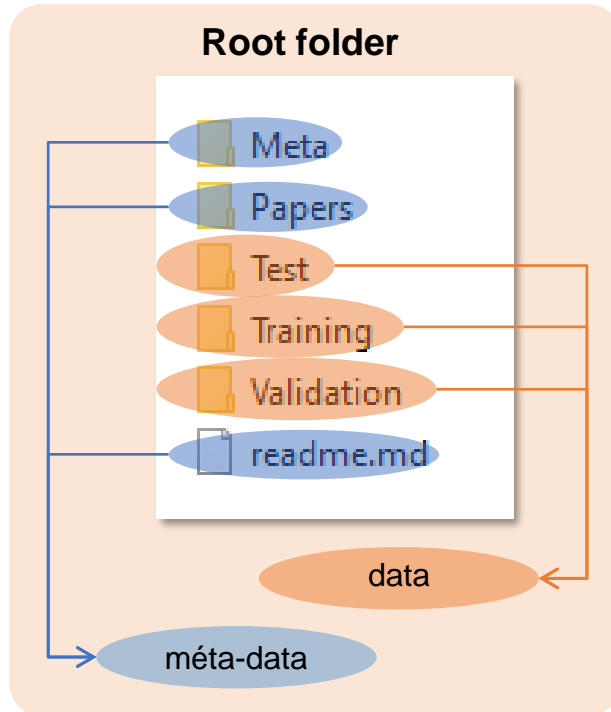
Prototype of **mobile application** that enables to take a picture of a fruit and get information on it.

→ Final goal: provide with a specific treatment for each fruit species through intelligent culling robots.



I. Introduction

2) The dataset




I. Introduction

3) Problematic


→ Currently 24 fruits in the database

→ Up to **650 photos per fruit**, in order to anticipate the different points of view of future images.




Problem: the database will grow with new fruits

→ risk of outgrowing regarding the **computing power** of a **single** computer.

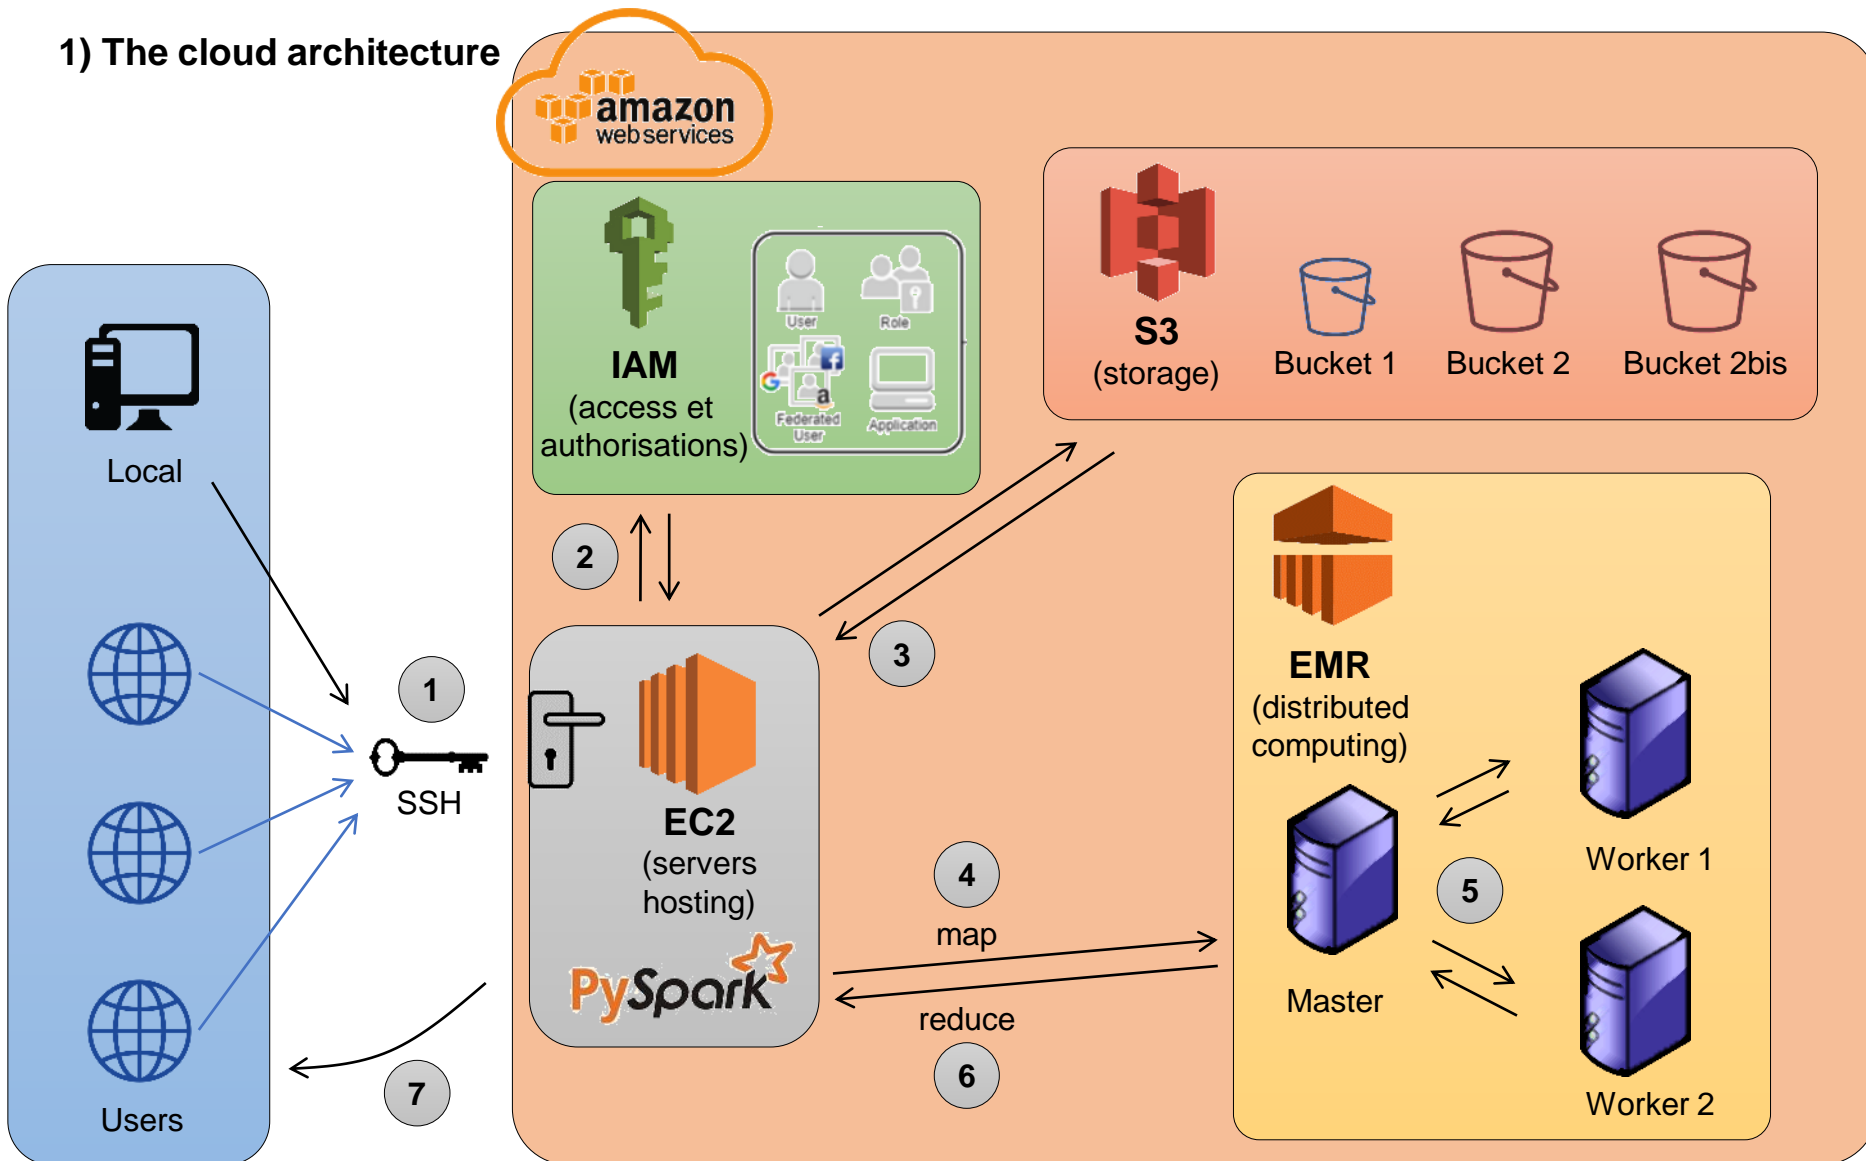


→ The future important data volume implies to scale up and use the **cloud** as well as **distributed computing**.



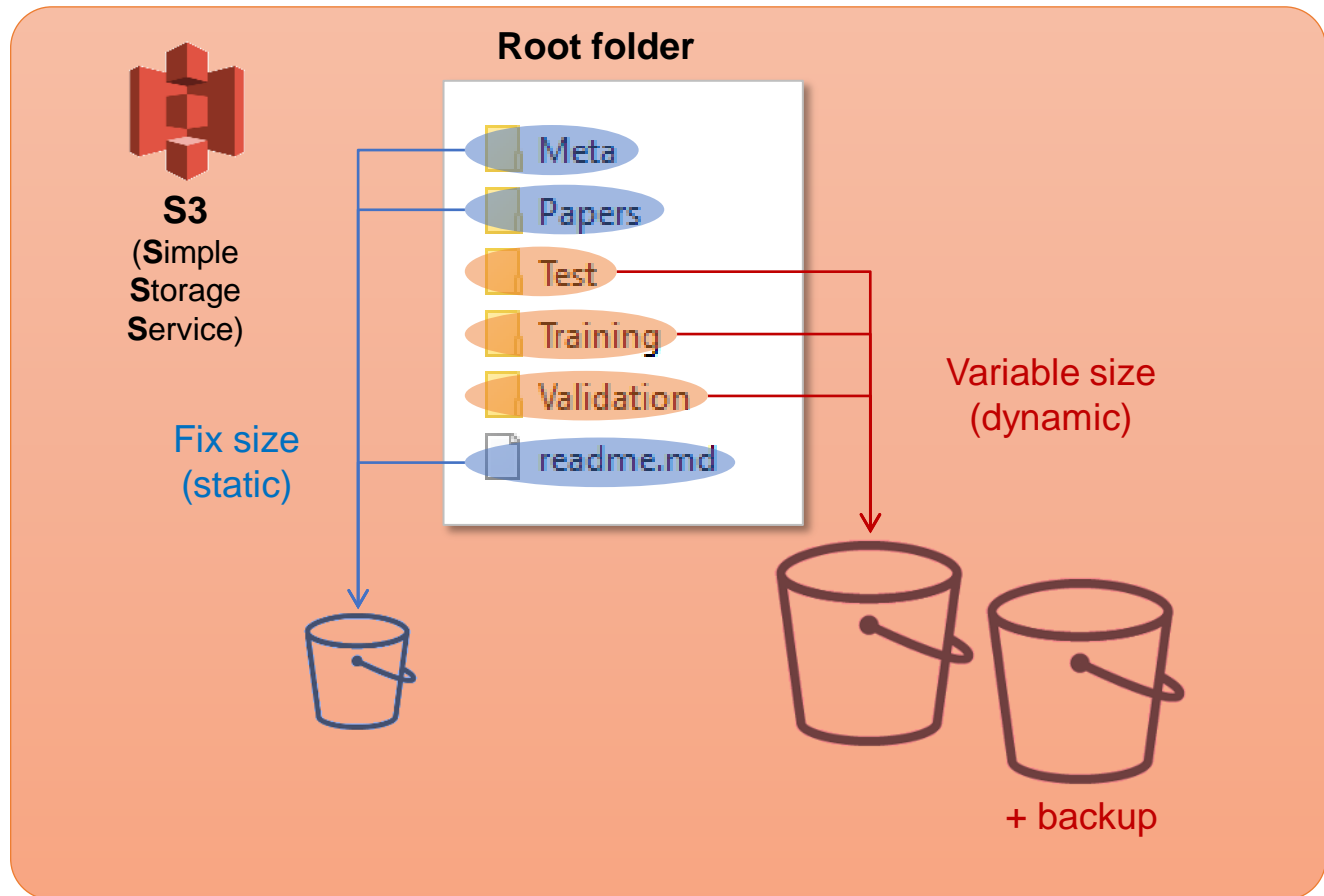
II. BigData environment

1) The cloud architecture



II. BigData environment

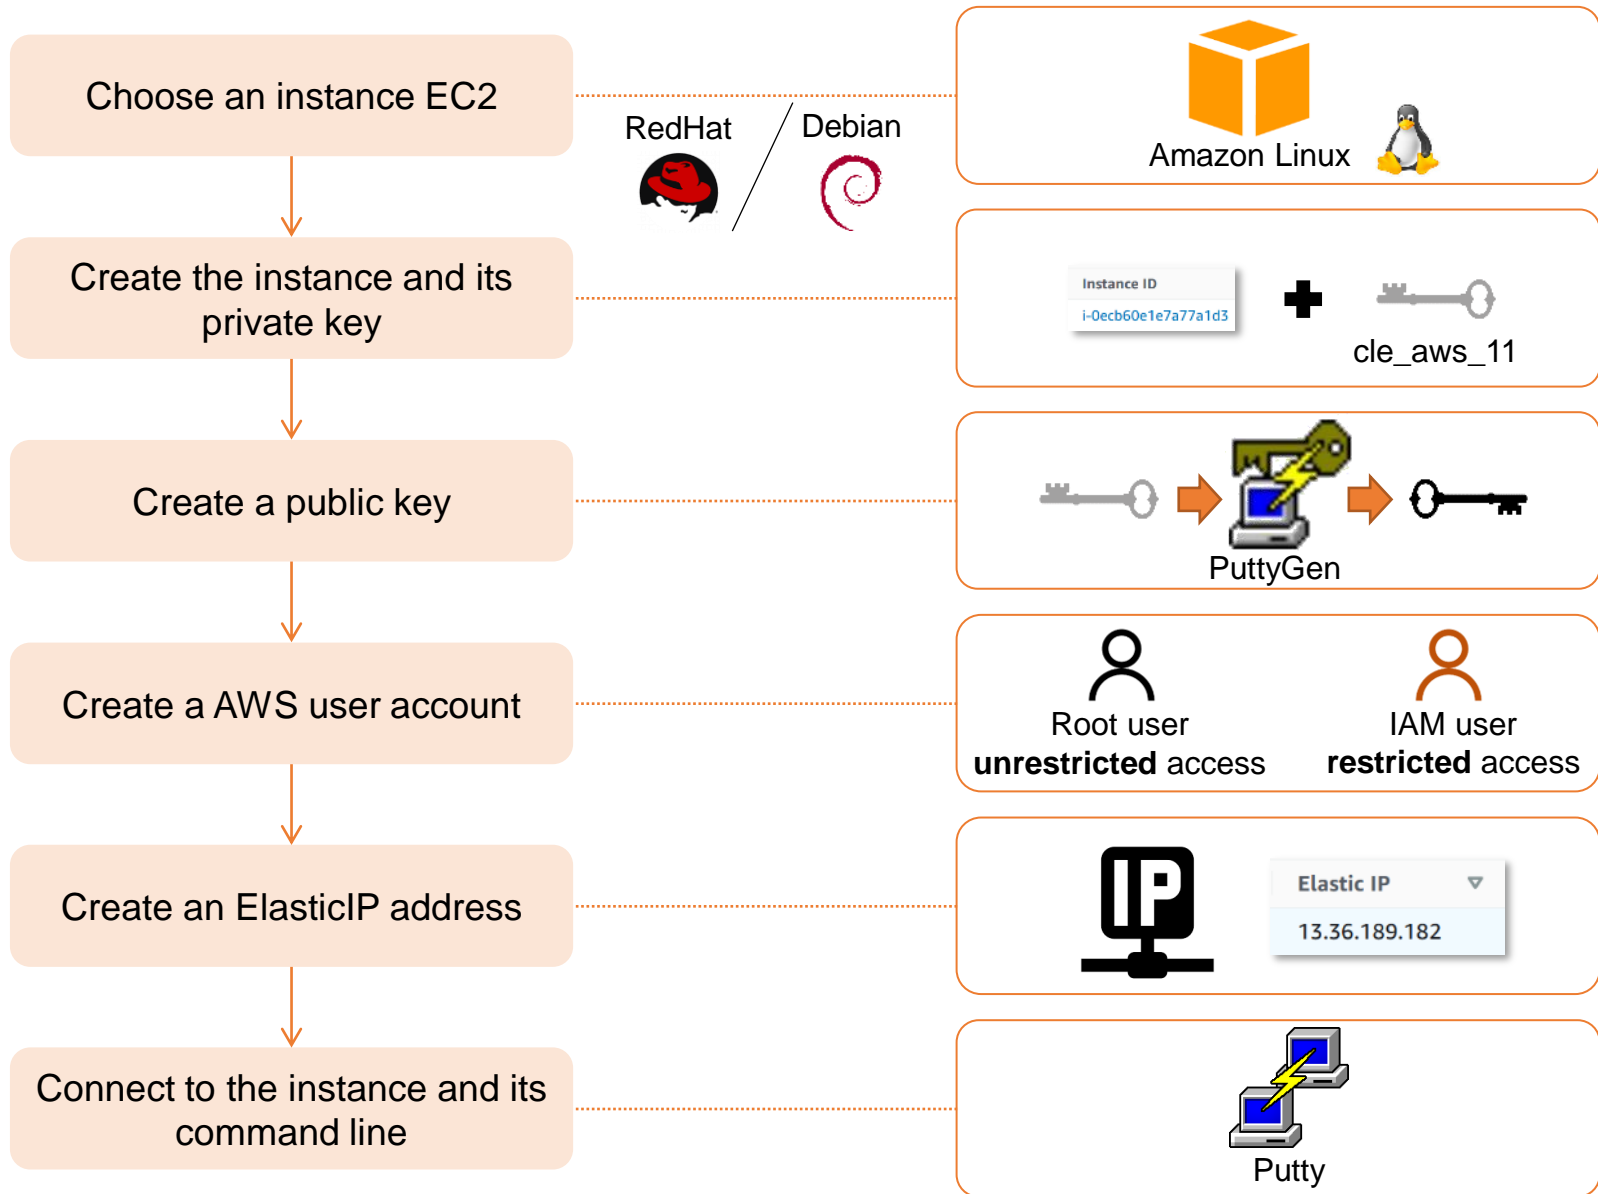
2) Data storage



<input type="radio"/>	ocr-taille-fixe	EU (Paris) eu-west-3	<u>Objects can be public</u>
<input type="radio"/>	ocr-taille-variable	EU (Paris) eu-west-3	<u>Objects can be public</u>
<input type="radio"/>	ocr-taille-variable-backup	EU (Paris) eu-west-3	<u>Bucket and objects not public</u>

II. BigData environment

3) Implementation of instance EC2



II. BigData environment

4) Preparation of Spark script

Implementation of security rules


1^{ère} try

Install Conda, Spark, .jar files

Install Tensorflow / Keras



Multiples errors of memory and version incompatibility



2^{ème} try

Choose an instance :

→ **t3.large** size

→ prepared for **deep learning**

→ with tensorflow et conda **pre-installed**

Type	Port range
Custom TCP	7077
SSH	22
Custom TCP	8443
Custom TCP	8888
HTTPS	443
Custom TCP	8080 - 8081

Activate Tensorflow & Conda

```
source activate tensorflow p37
```



Open a Jupyter notebook

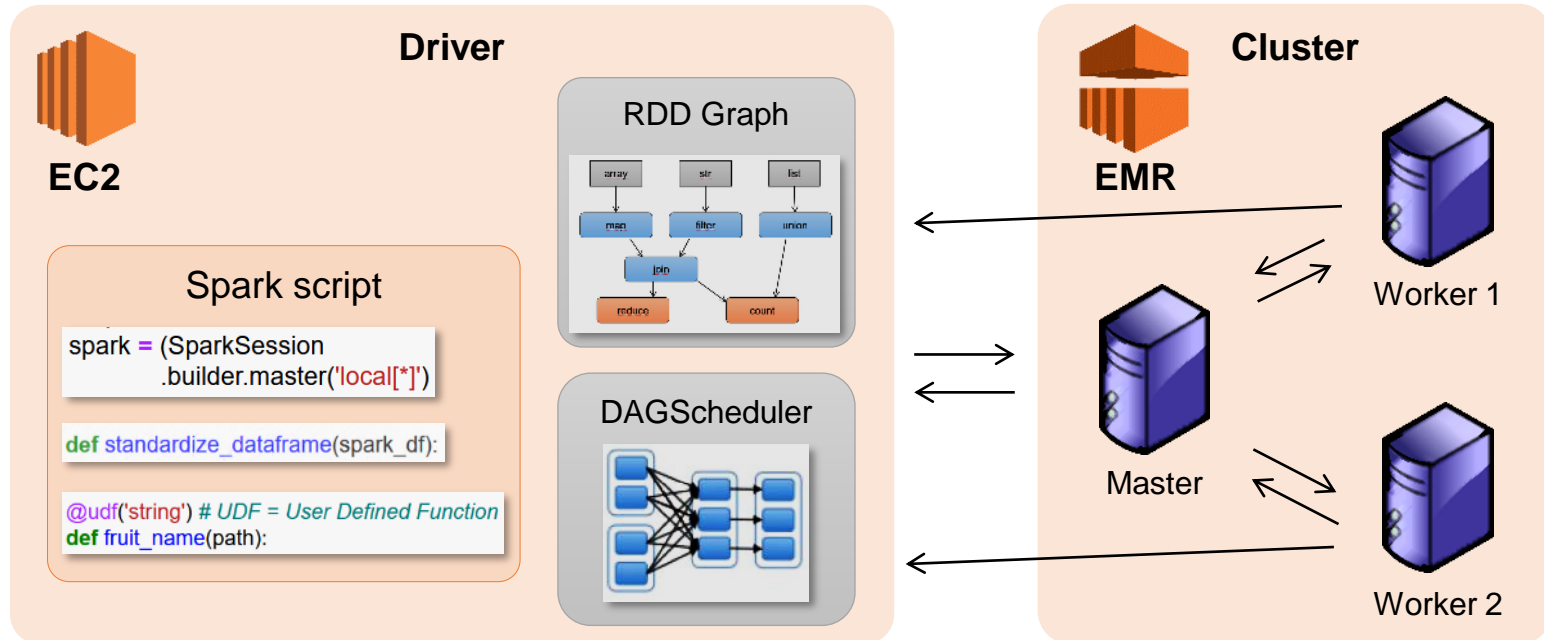


Write and execute the Spark script



II. BigData environment

5) Spark



Programme Spark

Execute Spark instructions

RDD Graph

Spark **Transformations** and **actions**

DAGScheduler

Distribute computations according Hadoop MapReduce scheme

Master

Shares computations between workers

Workers

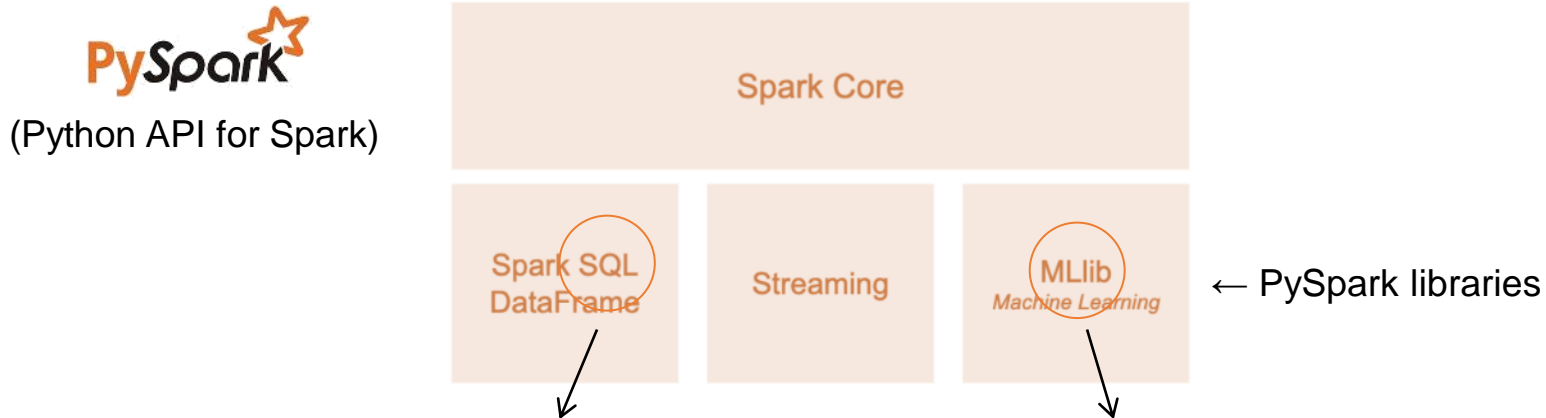
Execute computations

} Operations done by Spark

→ In the frame of this project, master and workers are hosted on the same server, i.e. the EC2 instance.

II. BigData environment

6) Libraries PySpark used



1) Methods of the library

```
from pyspark.sql import SparkSession  
  
spark = (SparkSession  
        .builder.master("local[*]"))
```

2) Decorators

```
from pyspark.sql.functions import udf  
  
@udf('string') # UDF = User Defined Function  
def fruit_name(path):
```

```
from pyspark.ml.feature import StandardScaler
```

```
standardizer = StandardScaler(withMean=True,  
                               withStd=True,  
                               inputCol='features_vector',  
                               outputCol='features_std')
```

→ Computations can be distributed with **methods of the PySpark library**, or **decorated functions**.

III. Image processing with Spark

2) Spark script configuration

→ A Spark script needs additional settings in comparison to an usual python script.

SparkSession


→ entry point of all functionality of Spark
→ encompasses all types of *contexts* :
Spark, Hive, SQL, ...

```
# Spark session
spark = SparkSession
    .builder.master('local[*]')
    .appName('p8_ocr')
    .config('spark.hadoop.fs.s3a.access.key', ACCESS_KEY_ID)
    .config('spark.hadoop.fs.s3a.secret.key', SECRET_ACCESS_KEY)
    .config('spark.hadoop.fs.s3a.impl', 'org.apache.hadoop.fs.s3a.S3AFileSystem')
    .config('com.amazonaws.services.s3.enableV4', 'true')
    .config('spark.hadoop.fs.s3a.endpoint', 's3.' + REGION + '.amazonaws.com')
    .getOrCreate()
```

SparkContext

→ connexion to a master-workers **cluster**


```
# Spark context and log level
spark_context = spark.sparkContext
spark_context.setLogLevel('WARN')
```



boto3

→ “*Python Software Development Kit*”
designed to configure and manage AWS
services.
→ helps to configure the connexion to **S3**

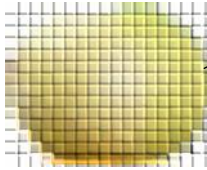
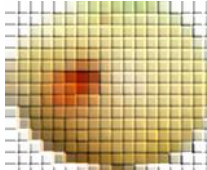
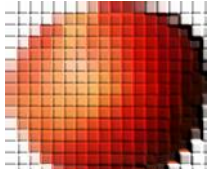
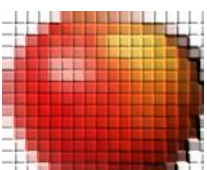
```
# S3 / EC2 authorizations
s3 = boto3.resource('s3',
                    REGION,
                    aws_access_key_id=ACCESS_KEY_ID,
                    aws_secret_access_key=SECRET_ACCESS_KEY)
bucket = s3.Bucket(BUCKET_NAME)
```



→ Now we can load images from S3 and use PySpark methods

III. Image processing with Spark

3) Raw table

fruit	jpeg	image
apple_6	r0_0.jpeg	
apple_6	r0_2.jpeg	
...
apple_braeburn_1	r0_0.jpeg	
apple_braeburn_1	r0_2.jpeg	
...

→ **Goal:** extract a table that can be used by an algorithm

III. Image processing with Spark

4.1) Image processing steps

Load images

```
def path_dataframe(img_list):
```

```
@udf('string') # UDF = User Defined Function  
def fruit_name(path):
```

```
@udf('string')  
def jpeg_name(path):
```

```
def readable_columns(path_sdf):
```

Create image table

Create readable columns

Predictions

```
resnet_model = ResNet50(include_top=False,  
                        weights=None,  
                        pooling='max',  
                        input_shape=(224, 224, 3))
```

Instantiate the neural network

```
def images_sample(image_list, img_nb=2400):
```

Select images randomly

```
def clean_image(image):
```

Load an image

```
def model_image_prediction(model, image):
```

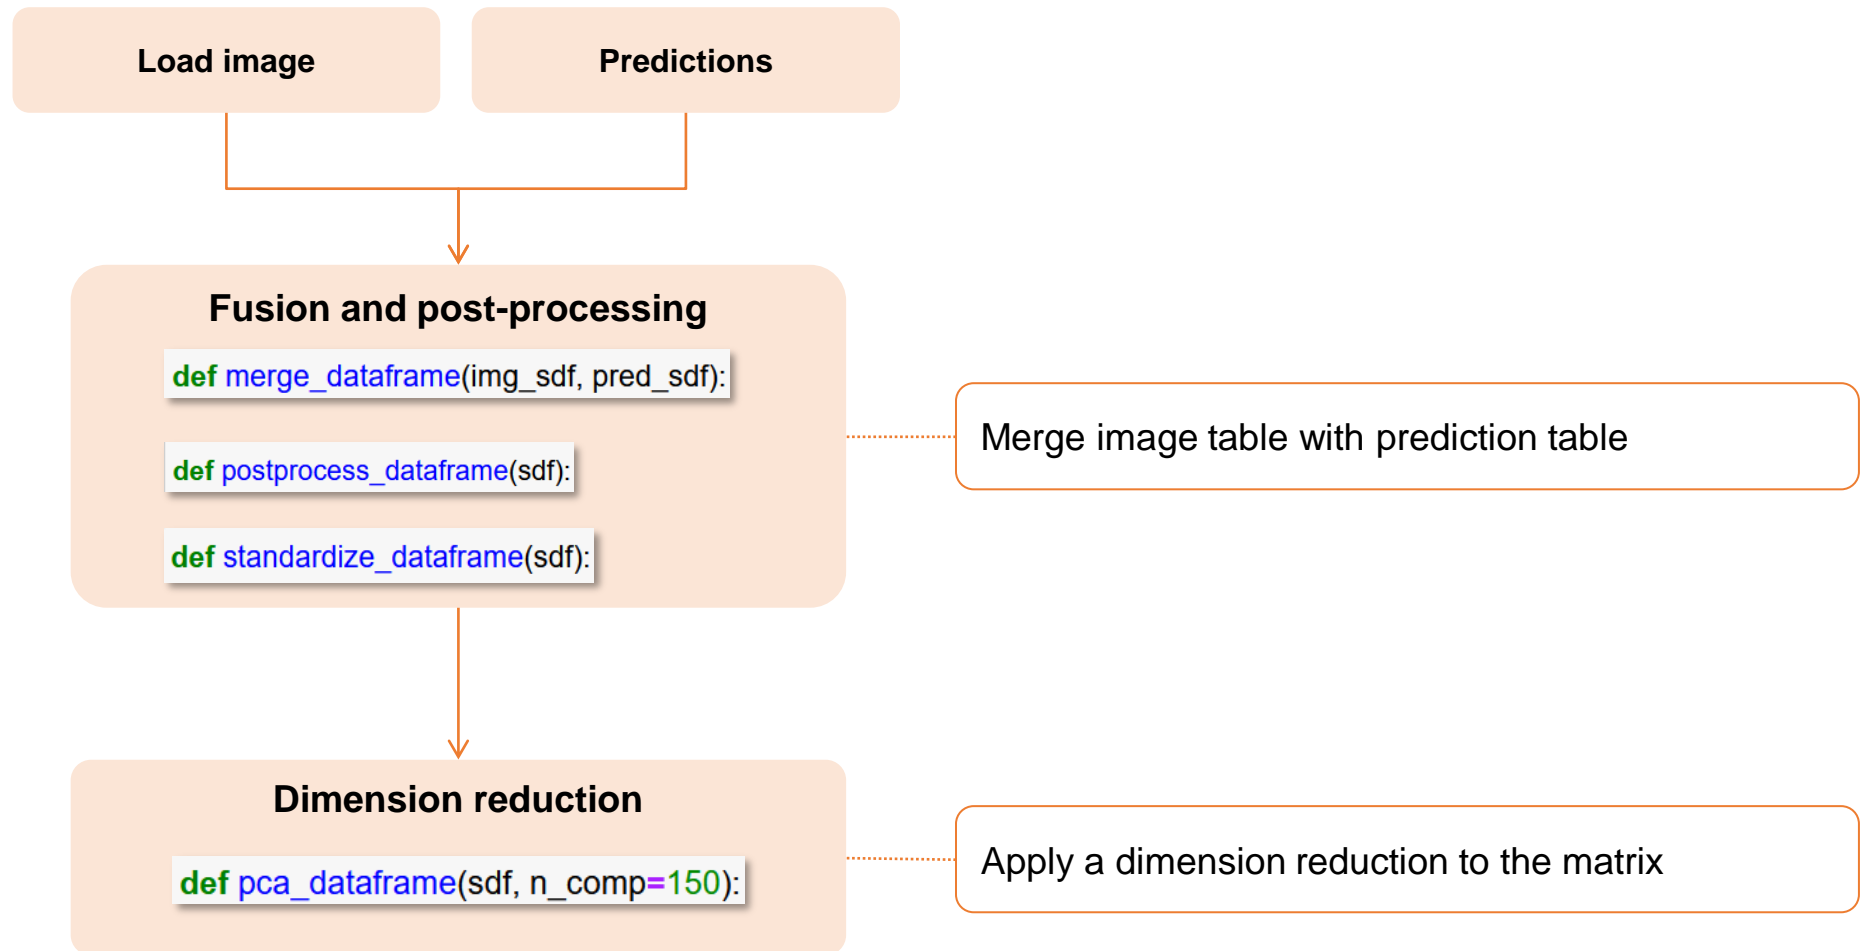
Get the prediction on an image

```
def predictions_dataframe(model, data, img_nb=200):
```

Gather all predictions

III. Image processing with Spark

4.2) Image processing steps



→ How many components for the reduced matrix ?


III. Image processing with Spark

5) PCA dimension reduction

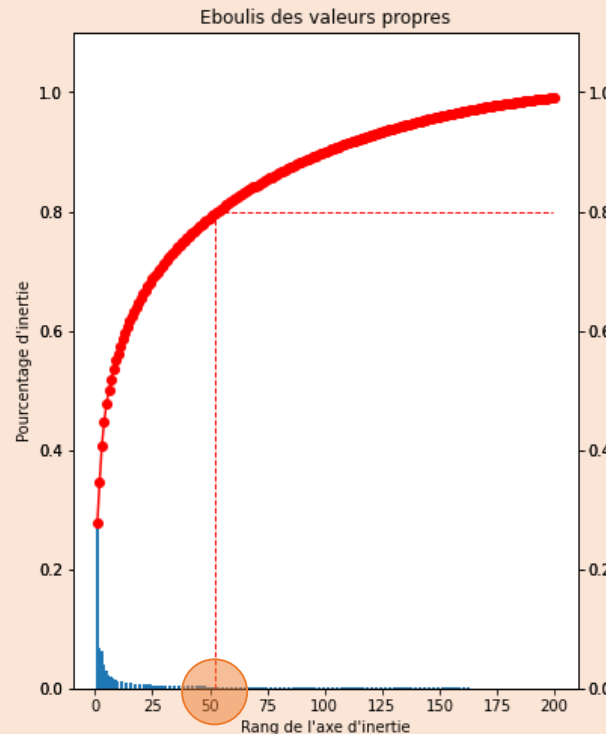
→ Weights of components vary with the number of images:

→ These graphs vary with the randomness of the `random.choices` method


→ Server utilization for 2400 images:



240 images (10 / fruit)
→ 80% in ≈ 50 components



2400 images (100 / fruit)
→ 80% in ≈ 140 components



→ As we increase the number of images processed, the 80% seem to converge to **150** principal components.

IV. Conclusion

1) Summary and outlooks



Dataset

- Ready for analysis in its current form.
- Growth is expected, so there is a need for **cloud** support (distributed computing)
- Majority of apple species (15 on 24 fruits): risk of bias.



BigData environnement

- Information exchanges take place between users, EC2 instance and S3 buckets.
- Instance **RAM** is important to be able to use Spark without trouble.
- **Version compatibility** is determining



Image processing with Spark

- Few images: need for a transfer learning with ResNet50 of Keras
- Distributed computations are done through **PySpark own methods** or **decorated functions**.
- PCA dimension reduction seems to be optimal with **150** components.



Outlooks

- Do parallelization of computations on a real EMR cluster.
- Diversify the fruits
- Consider other cloud solutions: Azure, Google Cloud, OVHCloud, ...
- Capture the images in equal proportions according to the rotation axes.
- Feature augmentation

IV. Conclusion

2) Recommendations for scaling

Quality criteria of a cloud tool

1. Simple scaling
2. Simple maintenance
3. Simple data exploitation

AWS 5 pillars

1. Operational excellence
2. Security
3. Reliability
4. Performance efficiency
5. Cost optimization

End of the presentation



Thank you for your attention